

lab[6]-report

57118123 刘康辉

Task 1: Implementing a Simple Firewall

用户主机的IP地址为10.9.0.5, 攻击者主机的IP地址为10.9.0.1, 内网主机的IP地址为192.168.60.5。

Task 1.A: Implement a Simple Kernel Module

创建文件夹, 保证完整路径中不含空格, 利用make命令编译可装载内核模块如下。

```
[07/22/21]seed@VM:~/.../Firewall$ make
make -C /lib/modules/5.4.0-54-generic/build M=/home/seed/Desktop/Firewall modules
make[1]: Entering directory '/usr/src/linux-headers-5.4.0-54-generic'
CC [M] /home/seed/Desktop/Firewall/hello.o
Building modules, stage 2.
MODPOST 1 modules
WARNING: modpost: missing MODULE_LICENSE() in /home/seed/Desktop/Firewall/hello.o
see include/linux/module.h for more information
CC [M] /home/seed/Desktop/Firewall/hello.mod.o
LD [M] /home/seed/Desktop/Firewall/hello.ko
make[1]: Leaving directory '/usr/src/linux-headers-5.4.0-54-generic'
```

利用insmod和rmmod命令将内核模块进行操作如下, 可知成功插入和移除。

```
[07/22/21]seed@VM:~/.../Firewall$ sudo insmod hello.ko
[07/22/21]seed@VM:~/.../Firewall$ lsmod | grep hello
hello                  16384  0
[07/22/21]seed@VM:~/.../Firewall$ sudo rmmod hello
```

利用dmesg命令查看/var/log/syslog文件中的信息, 得到结果如下。

```
[07/22/21]seed@VM:~/.../Firewall$ dmesg | grep World
[ 2494.132049] Hello World!
[ 2516.518487] Bye-bye World!.
```

Task 1.B: Implement a Simple Firewall Using Netfilter

1. Compile the sample code using the provided Makefile

在主机上利用dig查询www.example.com的DNS如下, 可知能够获得相关信息。

```
[07/22/21]seed@VM:~/.../Firewall$ dig @8.8.8.8 www.example.com

; <<>> DiG 9.16.1-Ubuntu <<>> @8.8.8.8 www.example.com
; (1 server found)
;; global options: +cmd
;; Got answer:
;; ->HEADER<<- opcode: QUERY, status: NOERROR, id: 47252
;; flags: qr rd ra ad; QUERY: 1, ANSWER: 1, AUTHORITY: 0, ADDITIONAL: 1

;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags:; udp: 512
;; QUESTION SECTION:
;www.example.com.                IN      A

;; ANSWER SECTION:
www.example.com.                18286   IN      A      93.184.216.34

;; Query time: 48 msec
;; SERVER: 8.8.8.8#53(8.8.8.8)
;; WHEN: Thu Jul 22 06:02:09 EDT 2021
;; MSG SIZE rcvd: 60
```

创建文件夹, 保证完整路径中不含空格, 利用make命令编译可装载内核模块如下。

```
[07/22/21]seed@VM:~/.../Firewall$ make
make -C /lib/modules/5.4.0-54-generic/build M=/home/seed/Desktop/Firewall modules
make[1]: Entering directory '/usr/src/linux-headers-5.4.0-54-generic'
CC [M] /home/seed/Desktop/Firewall/seedFilter.o
Building modules, stage 2.
MODPOST 1 modules
CC [M] /home/seed/Desktop/Firewall/seedFilter.mod.o
LD [M] /home/seed/Desktop/Firewall/seedFilter.ko
make[1]: Leaving directory '/usr/src/linux-headers-5.4.0-54-generic'
```

利用insmod命令插入内核模块，在主机上利用dig查询www.example.com的DNS如下，可知无法获得相关信息。

```
[07/22/21]seed@VM:~/.../Firewall$ sudo insmod seedFilter.ko
[07/22/21]seed@VM:~/.../Firewall$ dig @8.8.8.8 www.example.com
```

```
; <<>> DiG 9.16.1-Ubuntu <<>> @8.8.8.8 www.example.com
; (1 server found)
;; global options: +cmd
;; connection timed out; no servers could be reached
```

2. Hook the printInfo function to all of the netfilter hooks

修改seedFilter.c文件，代码如下。

```
#include <linux/kernel.h>
#include <linux/module.h>
#include <linux/netfilter.h>
#include <linux/netfilter_ipv4.h>
#include <linux/ip.h>
#include <linux/tcp.h>
#include <linux/udp.h>
#include <linux/if_ether.h>
#include <linux/inet.h>

static struct nf_hook_ops hook1, hook2, hook3, hook4, hook5;
unsigned int printInfo(void *priv, struct sk_buff *skb, const struct
nf_hook_state *state)
{
    struct iphdr *iph;
    char *hook;
    char *protocol;
    switch (state->hook){
        case NF_INET_LOCAL_IN:      hook = "LOCAL_IN";      break;
        case NF_INET_LOCAL_OUT:     hook = "LOCAL_OUT";     break;
        case NF_INET_PRE_ROUTING:   hook = "PRE_ROUTING";    break;
        case NF_INET_POST_ROUTING:  hook = "POST_ROUTING";   break;
        case NF_INET_FORWARD:       hook = "FORWARD";       break;
        default:                    hook = "IMPOSSIBLE";     break;
    }
    printk(KERN_INFO "**** %s\n", hook);
    iph = ip_hdr(skb);
    switch (iph->protocol){
        case IPPROTO_UDP:  protocol = "UDP";    break;
        case IPPROTO_TCP:  protocol = "TCP";    break;
        case IPPROTO_ICMP:  protocol = "ICMP";   break;
        default:            protocol = "OTHER";  break;
    }
    printk(KERN_INFO "      %pI4 --> %pI4 (%s)\n", &(iph->saddr), &(iph->daddr),
protocol);
    return NF_ACCEPT;
}
```

```

}
int registerFilter(void) {
    printk(KERN_INFO "Registering filters.\n");
    hook1.hook = printInfo;
    hook1.hooknum = NF_INET_PRE_ROUTING;
    hook1.pf = PF_INET;
    hook1.priority = NF_IP_PRI_FIRST;
    nf_register_net_hook(&init_net, &hook1);

    hook2.hook = printInfo;
    hook2.hooknum = NF_INET_LOCAL_IN;
    hook2.pf = PF_INET;
    hook2.priority = NF_IP_PRI_FIRST;
    nf_register_net_hook(&init_net, &hook2);

    hook3.hook = printInfo;
    hook3.hooknum = NF_INET_FORWARD;
    hook3.pf = PF_INET;
    hook3.priority = NF_IP_PRI_FIRST;
    nf_register_net_hook(&init_net, &hook3);

    hook4.hook = printInfo;
    hook4.hooknum = NF_INET_LOCAL_OUT;
    hook4.pf = PF_INET;
    hook4.priority = NF_IP_PRI_FIRST;
    nf_register_net_hook(&init_net, &hook4);

    hook5.hook = printInfo;
    hook5.hooknum = NF_INET_POST_ROUTING;
    hook5.pf = PF_INET;
    hook5.priority = NF_IP_PRI_FIRST;
    nf_register_net_hook(&init_net, &hook5);
    return 0;
}
void removeFilter(void) {
    printk(KERN_INFO "The filters are being removed.\n");
    nf_unregister_net_hook(&init_net, &hook1);
    nf_unregister_net_hook(&init_net, &hook2);
}
module_init(registerFilter);
module_exit(removeFilter);
MODULE_LICENSE("GPL");

```

利用make命令编译可装载内核模块，并且利用insmod命令插入内核模块如下。

```

[07/23/21]seed@VM:~/.../Firewall$ make
make -C /lib/modules/5.4.0-54-generic/build M=/home/seed/Desktop/Firewall modules
make[1]: Entering directory '/usr/src/linux-headers-5.4.0-54-generic'
  CC [M]  /home/seed/Desktop/Firewall/seedFilter.o
Building modules, stage 2.
MODPOST 1 modules
  CC [M]  /home/seed/Desktop/Firewall/seedFilter.mod.o
  LD [M]  /home/seed/Desktop/Firewall/seedFilter.ko
make[1]: Leaving directory '/usr/src/linux-headers-5.4.0-54-generic'
[07/23/21]seed@VM:~/.../Firewall$ sudo insmod seedFilter.ko
[07/23/21]seed@VM:~/.../Firewall$ lsmod | grep seedFilter
seedFilter                16384  0

```

在用户主机上ping内网主机，得到结果如下，可知能够连接。

```
root@36cddb526b95:/# ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
64 bytes from 192.168.60.5: icmp_seq=1 ttl=63 time=0.222 ms
```

利用dmesg命令查看/var/log/syslog文件中的信息，得到结果如下。

```
[ 469.983357] *** PRE_ROUTING
[ 469.983361]      10.9.0.5 --> 192.168.60.5 (ICMP)
[ 469.983372] *** FORWARD
[ 469.983373]      10.9.0.5 --> 192.168.60.5 (ICMP)
[ 469.983377] *** POST_ROUTING
[ 469.983378]      10.9.0.5 --> 192.168.60.5 (ICMP)
```

在用户主机上ping攻击者主机，得到结果如下，可知能够连接。

```
root@36cddb526b95:/# ping 10.9.0.1
PING 10.9.0.1 (10.9.0.1) 56(84) bytes of data.
64 bytes from 10.9.0.1: icmp_seq=1 ttl=64 time=0.132 ms
```

利用dmesg命令查看/var/log/syslog文件中的信息，得到结果如下。

```
[ 704.883109] *** PRE_ROUTING
[ 704.883110]      10.9.0.5 --> 10.9.0.1 (ICMP)
[ 704.883113] *** LOCAL_IN
[ 704.883114]      10.9.0.5 --> 10.9.0.1 (ICMP)
[ 704.883123] *** LOCAL_OUT
[ 704.883124]      10.9.0.1 --> 10.9.0.5 (ICMP)
[ 704.883126] *** POST_ROUTING
[ 704.883127]      10.9.0.1 --> 10.9.0.5 (ICMP)
```

根据实验结果可知，NF_IP_PRE_ROUTING在数据包刚进入主机进行处理的时候调用；NF_IP_LOCAL_IN在确认数据包的目的地址为本机的时候调用；NF_IP_FORWARD在要数据包通过主机进行转发的时候调用；NF_IP_LOCAL_OUT在确认数据包的源地址为本机的时候调用；NF_IP_POST_ROUTING在数据包将离开主机进行处理的时候调用。

3. Implement two more hooks

修改seedFilter.c文件，代码如下。

```
#include <linux/kernel.h>
#include <linux/module.h>
#include <linux/netfilter.h>
#include <linux/netfilter_ipv4.h>
#include <linux/ip.h>
#include <linux/tcp.h>
#include <linux/udp.h>
#include <linux/if_ether.h>
#include <linux/inet.h>

static struct nf_hook_ops hook1, hook2;
unsigned int ICMPFilter(void *priv, struct sk_buff *skb, const struct
nf_hook_state *state)
{
    struct iphdr *iph;
    iph = ip_hdr(skb);

    if (iph->protocol == IPPROTO_ICMP) {
        printk(KERN_INFO "Dropping ICMP packet:%pI4\n", &(iph->saddr));
        return NF_DROP;
    }
    return NF_ACCEPT;
}
```

```

}
unsigned int telnetFilter(void *priv, struct sk_buff *skb, const struct
nf_hook_state *state)
{
    struct iphdr *iph;
    struct tcphdr *tcph;
    iph = ip_hdr(skb);
    tcph = (void *)iph + iph->ihl * 4;
    if (iph->protocol == IPPROTO_TCP && tcph->dest == htons(23)) {
        printk(KERN_INFO "Dropping telnet packet:%pI4\n", &(iph->saddr));
        return NF_DROP;
    }
    return NF_ACCEPT;
}
int registerFilter(void) {
    printk(KERN_INFO "Registering filters.\n");
    hook1.hook = ICMPFilter;
    hook1.hooknum = NF_INET_LOCAL_IN;
    hook1.pf = PF_INET;
    hook1.priority = NF_IP_PRI_FIRST;
    nf_register_net_hook(&init_net, &hook1);

    hook2.hook = telnetFilter;
    hook2.hooknum = NF_INET_LOCAL_IN;
    hook2.pf = PF_INET;
    hook2.priority = NF_IP_PRI_FIRST;
    nf_register_net_hook(&init_net, &hook2);
    return 0;
}
void removeFilter(void) {
    printk(KERN_INFO "The filters are being removed.\n");
    nf_unregister_net_hook(&init_net, &hook1);
    nf_unregister_net_hook(&init_net, &hook2);
}
module_init(registerFilter);
module_exit(removeFilter);
MODULE_LICENSE("GPL");

```

利用make命令编译可装载内核模块，并且利用insmod命令插入内核模块如下。

```

[07/23/21]seed@VM:~/.../Firewall$ make
make -C /lib/modules/5.4.0-54-generic/build M=/home/seed/Desktop/Firewall modules
make[1]: Entering directory '/usr/src/linux-headers-5.4.0-54-generic'
  CC [M] /home/seed/Desktop/Firewall/seedFilter.o
  Building modules, stage 2.
  MODPOST 1 modules
  CC [M] /home/seed/Desktop/Firewall/seedFilter.mod.o
  LD [M] /home/seed/Desktop/Firewall/seedFilter.ko
make[1]: Leaving directory '/usr/src/linux-headers-5.4.0-54-generic'
[07/23/21]seed@VM:~/.../Firewall$ sudo insmod seedFilter.ko
[07/23/21]seed@VM:~/.../Firewall$ lsmod | grep seedFilter
seedFilter                16384  0

```

在用户主机上ping攻击者主机，得到结果如下，可知无法连接。

```

root@53fd31c6c3d3:/# ping 10.9.0.1
PING 10.9.0.1 (10.9.0.1) 56(84) bytes of data.
^C
--- 10.9.0.1 ping statistics ---
14 packets transmitted, 0 received, 100% packet loss, time 13318ms

```

在用户主机上telnet远程连接攻击者主机，得到结果如下，可知连接失败。

```
root@53fd31c6c3d3:/# telnet 10.9.0.1
Trying 10.9.0.1...
telnet: Unable to connect to remote host: Connection timed out
```

利用dmesg命令查看/var/log/syslog文件中的信息，得到结果如下，可知ICMP和telnet报文都被丢弃。

```
[07/23/21]seed@VM:~/.../Firewall$ dmesg | grep Dropping
[ 7508.610373] Dropping ICMP packet:10.9.0.5
[ 7509.633199] Dropping ICMP packet:10.9.0.5
[ 7510.654358] Dropping ICMP packet:10.9.0.5
[ 7511.679002] Dropping ICMP packet:10.9.0.5
[ 7512.705140] Dropping ICMP packet:10.9.0.5
[ 7513.727040] Dropping ICMP packet:10.9.0.5
[ 7514.752881] Dropping ICMP packet:10.9.0.5
[ 7515.775368] Dropping ICMP packet:10.9.0.5
[ 7516.800981] Dropping ICMP packet:10.9.0.5
[ 7517.838316] Dropping ICMP packet:10.9.0.5
[ 7518.858580] Dropping ICMP packet:10.9.0.5
[ 7519.878769] Dropping ICMP packet:10.9.0.5
[ 7520.898500] Dropping ICMP packet:10.9.0.5
[ 7521.928230] Dropping ICMP packet:10.9.0.5
[ 7529.024203] Dropping telnet packet:10.9.0.5
[ 7530.086482] Dropping telnet packet:10.9.0.5
[ 7532.320529] Dropping telnet packet:10.9.0.5
[ 7536.415275] Dropping telnet packet:10.9.0.5
[ 7544.607365] Dropping telnet packet:10.9.0.5
[ 7560.734450] Dropping telnet packet:10.9.0.5
[ 7592.992919] Dropping telnet packet:10.9.0.5
```

Task 2: Experimenting with Stateless Firewall Rules

用户主机的IP地址为10.9.0.5，路由器的IP地址为10.9.0.11，内网网段的IP地址为192.168.60.0/24。

Task 2.A: Protecting the Router

在路由器上利用iptables命令，创建过滤规则如下。

```
root@53fc248d8be9:/# iptables -A INPUT -p icmp --icmp-type echo-request -j ACCEPT
root@53fc248d8be9:/# iptables -A OUTPUT -p icmp --icmp-type echo-reply -j ACCEPT
root@53fc248d8be9:/# iptables -P OUTPUT DROP
root@53fc248d8be9:/# iptables -P INPUT DROP
```

在用户主机上ping路由器，得到结果如下，可知能够连接。

```
root@53fde8986d54:/# ping 10.9.0.11
PING 10.9.0.11 (10.9.0.11) 56(84) bytes of data.
64 bytes from 10.9.0.11: icmp_seq=1 ttl=64 time=0.105 ms
64 bytes from 10.9.0.11: icmp_seq=2 ttl=64 time=0.060 ms
64 bytes from 10.9.0.11: icmp_seq=3 ttl=64 time=0.109 ms
```

在用户主机上telnet远程连接路由器，得到结果如下，可知连接失败。

```
root@53fde8986d54:/# telnet 10.9.0.11
Trying 10.9.0.11...
telnet: Unable to connect to remote host: Connection timed out
```

该现象的原因是路由器的过滤规则只允许icmp请求报文输入和icmp响应报文输出，ping的报文可以进行传输，而telnet的报文无法进行传输。

Task 2.B: Protecting the Internal Network

在路由器上利用iptables命令，创建过滤规则如下。

```
root@53fc248d8be9:/# iptables -A FORWARD -p icmp --icmp-type echo-request -i eth1 -j ACCEPT
root@53fc248d8be9:/# iptables -A FORWARD -p icmp --icmp-type echo-reply -i eth0 -j ACCEPT
root@53fc248d8be9:/# iptables -P FORWARD DROP
```

在用户主机上ping内网主机192.168.60.5，得到结果如下，可知无法连接。


```
root@53fde8986d54:/# ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
^C
--- 192.168.60.5 ping statistics ---
4 packets transmitted, 0 received, 100% packet loss, time 3061ms
```

在用户主机上ping路由器，得到结果如下，可知能够连接。

```
root@53fde8986d54:/# ping 10.9.0.11
PING 10.9.0.11 (10.9.0.11) 56(84) bytes of data.
64 bytes from 10.9.0.11: icmp_seq=1 ttl=64 time=0.167 ms
64 bytes from 10.9.0.11: icmp_seq=2 ttl=64 time=0.087 ms
64 bytes from 10.9.0.11: icmp_seq=3 ttl=64 time=0.073 ms
```

在IP地址为192.168.60.5的内网主机上ping用户主机，得到结果如下，可知能够连接。

```
root@df9d5cbeb2c8:/# ping 10.9.0.5
PING 10.9.0.5 (10.9.0.5) 56(84) bytes of data.
64 bytes from 10.9.0.5: icmp_seq=1 ttl=63 time=0.168 ms
64 bytes from 10.9.0.5: icmp_seq=2 ttl=63 time=0.108 ms
64 bytes from 10.9.0.5: icmp_seq=3 ttl=63 time=0.062 ms
```

在用户主机上telnet远程连接内网主机192.168.60.5，得到结果如下，可知连接失败。

```
root@53fde8986d54:/# telnet 192.168.60.5
Trying 192.168.60.5...
telnet: Unable to connect to remote host: Connection timed out
```

在IP地址为192.168.60.5的内网主机上telnet远程连接用户主机，得到结果如下，可知连接失败。

```
root@df9d5cbeb2c8:/# telnet 10.9.0.5
Trying 10.9.0.5...
telnet: Unable to connect to remote host: Connection timed out
```

Task 2.C: Protecting Internal Servers

在路由器上利用iptables命令，创建过滤规则如下。

```
root@53fc248d8be9:/# iptables -A FORWARD -i eth0 -p tcp -d 192.168.60.5 --dport 23 -j ACCEPT
root@53fc248d8be9:/# iptables -A FORWARD -o eth0 -p tcp -s 192.168.60.5 --sport 23 -j ACCEPT
root@53fc248d8be9:/# iptables -A FORWARD -i eth1 -p tcp -d 192.168.60.0/24 --dport 23 -j ACCEPT
root@53fc248d8be9:/# iptables -A FORWARD -o eth1 -p tcp -s 192.168.60.0/24 --sport 23 -j ACCEPT
root@53fc248d8be9:/# iptables -P FORWARD DROP
```

在用户主机上telnet远程连接内网主机192.168.60.5，得到结果如下，可知连接成功。

```
root@53fde8986d54:/# telnet 192.168.60.5
Trying 192.168.60.5...
Connected to 192.168.60.5.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
df9d5cbeb2c8 login: seed
Password:
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.4.0-54-generic x86_64)
```

在用户主机上telnet远程连接内网主机192.168.60.6，得到结果如下，可知连接失败。

```
root@53fde8986d54:/# telnet 192.168.60.6
Trying 192.168.60.6...
telnet: Unable to connect to remote host: Connection timed out
```

在IP地址为192.168.60.5的内网主机上telnet远程连接内网主机192.168.60.6，得到结果如下，可知连接成功。

```
root@df9d5cbeb2c8:/# telnet 192.168.60.6
Trying 192.168.60.6...
Connected to 192.168.60.6.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
c15da65f559c login: seed
Password:
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.4.0-54-generic x86_64)
```

在IP地址为192.168.60.5的内网主机上telnet远程连接用户主机，得到结果如下，可知连接失败。

```
root@df9d5cbeb2c8:/# telnet 10.9.0.5
Trying 10.9.0.5...
telnet: Unable to connect to remote host: Connection timed out
```

Task 3: Connection Tracking and Stateful Firewall

用户主机的IP地址为10.9.0.5，路由器的IP地址为10.9.0.11，内网网段的IP地址为192.168.60.0/24。

Task 3.A: Experiment with the Connection Tracking

ICMP experiment

在用户主机上ping内网主机192.168.60.5，得到结果如下，可知能够连接。

```
root@53fde8986d54:/# ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
64 bytes from 192.168.60.5: icmp_seq=1 ttl=63 time=0.250 ms
64 bytes from 192.168.60.5: icmp_seq=2 ttl=63 time=0.073 ms
64 bytes from 192.168.60.5: icmp_seq=3 ttl=63 time=0.115 ms
```

在路由器上利用conntrack -L命令实现连接跟踪，得到结果如下，可知ICMP连接时间约为30s。

```
root@53fc248d8be9:/# conntrack -L
icmp      1 29 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=57 src=192.168.60.5
dst=10.9.0.5 type=0 code=0 id=57 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
```

UDP experiment

在用户主机上利用UDP远程连接IP地址为192.168.60.5的内网主机9090端口，并发送消息如下。

```
root@53fde8986d54:/# nc -u 192.168.60.5 9090
whoami
```

在内网主机192.168.60.5上监听9090端口的UDP连接，得到结果如下。

```
root@df9d5cbeb2c8:/# nc -lu 9090
whoami
```

在路由器上利用conntrack -L命令实现连接跟踪，得到结果如下，可知UDP连接时间约为30s。

```
root@53fc248d8be9:/# conntrack -L
udp      17 28 src=10.9.0.5 dst=192.168.60.5 sport=49265 dport=9090 [UNREPLIED]
src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=49265 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
```

TCP experiment

在用户主机上利用TCP远程连接IP地址为192.168.60.5的内网主机9090端口，并发送消息如下。

```
root@53fde8986d54:/# nc 192.168.60.5 9090
whoami
```

在内网主机192.168.60.5上监听9090端口的TCP连接，得到结果如下。


```
root@df9d5cbeb2c8:/# nc -l 9090
whoami
```

在路由器上利用conntrack -L命令实现连接跟踪，得到结果如下，可知TCP连接时间约为432000s。

```
root@53fc248d8be9:/# conntrack -L
tcp        6 431997 ESTABLISHED src=10.9.0.5 dst=192.168.60.5 sport=41558 dport=90
90 src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=41558 [ASSURED] mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
```

Task 3.B: Setting Up a Stateful Firewall

在路由器上利用iptables命令和连接跟踪机制，创建过滤规则如下。

```
root@53fc248d8be9:/# iptables -A FORWARD -p tcp -m conntrack --ctstate ESTABLISHED,RELATED -j ACCEPT
root@53fc248d8be9:/# iptables -A FORWARD -p tcp -i eth0 -d 192.168.60.5 --dport 23 --syn -m conntrack --ctstate NEW -j ACCEPT
root@53fc248d8be9:/# iptables -A FORWARD -p tcp -i eth1 -s 192.168.60.0/24 --dport 23 --syn -m conntrack --ctstate NEW -j ACCEPT
root@53fc248d8be9:/# iptables -P FORWARD DROP
```

在用户主机上telnet远程连接内网主机192.168.60.5，得到结果如下，可知连接成功。

```
root@53fde8986d54:/# telnet 192.168.60.5
Trying 192.168.60.5...
Connected to 192.168.60.5.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
df9d5cbeb2c8 login: seed
Password:
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.4.0-54-generic x86_64)
```

在用户主机上telnet远程连接内网主机192.168.60.6，得到结果如下，可知连接失败。

```
root@53fde8986d54:/# telnet 192.168.60.6
Trying 192.168.60.6...
telnet: Unable to connect to remote host: Connection timed out
```

在IP地址为192.168.60.5的内网主机上telnet远程连接内网主机192.168.60.6，得到结果如下，可知连接成功。

```
root@df9d5cbeb2c8:/# telnet 192.168.60.6
Trying 192.168.60.6...
Connected to 192.168.60.6.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
c15da65f559c login: seed
Password:
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.4.0-54-generic x86_64)
```

在IP地址为192.168.60.5的内网主机上telnet远程连接用户主机，得到结果如下，可知连接成功。

```
root@df9d5cbeb2c8:/# telnet 10.9.0.5
Trying 10.9.0.5...
Connected to 10.9.0.5.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
53fde8986d54 login: seed
Password:
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.4.0-54-generic x86_64)
```

不利用连接跟踪机制的过滤规则仅对数据包的首部进行检查，其优点是处理速度快，缺点是无法定义精细的规则、不适合复杂的访问控制；而利用连接跟踪机制的过滤规则对数据包的状态也进行检查，其优点是能够定义更加严格的规则、应用范围更广、安全性更高，缺点是无法对数据包的内容进行识别。

Task 4: Limiting Network Traffic

在路由器上利用iptables命令，创建流量限制规则如下。

```
root@53fc248d8be9:/# iptables -A FORWARD -s 10.9.0.5 -m limit --limit 10/minute
--limit-burst 5 -j ACCEPT
root@53fc248d8be9:/# iptables -A FORWARD -s 10.9.0.5 -j DROP
```

在用户主机上ping内网主机192.168.60.5，得到结果如下，可知能够连接，但部分报文因流量限制而丢失。

```
root@53fde8986d54:/# ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
64 bytes from 192.168.60.5: icmp_seq=1 ttl=63 time=0.099 ms
64 bytes from 192.168.60.5: icmp_seq=2 ttl=63 time=0.082 ms
64 bytes from 192.168.60.5: icmp_seq=3 ttl=63 time=0.112 ms
64 bytes from 192.168.60.5: icmp_seq=4 ttl=63 time=0.111 ms
64 bytes from 192.168.60.5: icmp_seq=5 ttl=63 time=0.075 ms
64 bytes from 192.168.60.5: icmp_seq=7 ttl=63 time=0.076 ms
64 bytes from 192.168.60.5: icmp_seq=13 ttl=63 time=0.093 ms
^C
--- 192.168.60.5 ping statistics ---
14 packets transmitted, 7 received, 50% packet loss, time 13313ms
rtt min/avg/max/mdev = 0.075/0.092/0.112/0.014 ms
```

在路由器上利用iptables命令，修改流量限制规则如下。

```
root@53fc248d8be9:/# iptables -A FORWARD -s 10.9.0.5 -m limit --limit 10/minute
--limit-burst 5 -j ACCEPT
```

在用户主机上ping内网主机192.168.60.5，得到结果如下，可知能够连接，且无报文丢失。

```
root@53fde8986d54:/# ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
64 bytes from 192.168.60.5: icmp_seq=1 ttl=63 time=0.116 ms
64 bytes from 192.168.60.5: icmp_seq=2 ttl=63 time=0.086 ms
64 bytes from 192.168.60.5: icmp_seq=3 ttl=63 time=0.119 ms
64 bytes from 192.168.60.5: icmp_seq=4 ttl=63 time=0.077 ms
64 bytes from 192.168.60.5: icmp_seq=5 ttl=63 time=0.076 ms
64 bytes from 192.168.60.5: icmp_seq=6 ttl=63 time=0.099 ms
64 bytes from 192.168.60.5: icmp_seq=7 ttl=63 time=0.106 ms
64 bytes from 192.168.60.5: icmp_seq=8 ttl=63 time=0.106 ms
64 bytes from 192.168.60.5: icmp_seq=9 ttl=63 time=0.079 ms
64 bytes from 192.168.60.5: icmp_seq=10 ttl=63 time=0.108 ms
64 bytes from 192.168.60.5: icmp_seq=11 ttl=63 time=0.096 ms
64 bytes from 192.168.60.5: icmp_seq=12 ttl=63 time=0.108 ms
64 bytes from 192.168.60.5: icmp_seq=13 ttl=63 time=0.110 ms
^C
--- 192.168.60.5 ping statistics ---
13 packets transmitted, 13 received, 0% packet loss, time 12293ms
rtt min/avg/max/mdev = 0.076/0.098/0.119/0.014 ms
```

该现象的原因是路由器的转发链的默认规则为ACCEPT，即使超过流量限制，报文根据默认规则也可以进行传输，可知上述第二条规则是必需的。

Task 5: Load Balancing

用户主机的IP地址为10.9.0.5，路由器的IP地址为10.9.0.11，三个服务器的IP地址为192.168.60.5、192.168.60.6和192.168.60.7。

Using the nth mode (round-robin)

在路由器上利用iptables命令，采用nth模式创建负载均衡规则如下。

```
root@eaf163c47638:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic
--mode nth --every 3 --packet 0 -j DNAT --to-destination 192.168.60.5:8080
root@eaf163c47638:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic
--mode nth --every 2 --packet 0 -j DNAT --to-destination 192.168.60.6:8080
root@eaf163c47638:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -j DNAT --to-
destination 192.168.60.7:8080
```

在用户主机上向路由器的8080端口发送UDP数据包如下。

```
root@36cddb526b95:/# echo hello | nc -u 10.9.0.11 8080
^C
root@36cddb526b95:/# echo hello | nc -u 10.9.0.11 8080
^C
root@36cddb526b95:/# echo hello | nc -u 10.9.0.11 8080
^C
root@36cddb526b95:/# echo hello | nc -u 10.9.0.11 8080
^C
```

在服务器192.168.60.5上监听8080端口，得到结果如下。

```
root@44ca35555561:/# nc -luk 8080
hello
hello
```

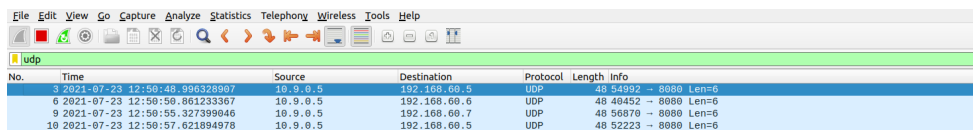
在服务器192.168.60.6上监听8080端口，得到结果如下。

```
root@5479b7208878:/# nc -luk 8080
hello
```

在服务器192.168.60.7上监听8080端口，得到结果如下。

```
root@df14eb7bfad1:/# nc -luk 8080
hello
```

利用wireshark抓包，得到结果如下，可知UDP数据包发送成功，且根据负载均衡规则依次发给三个服务器。



No.	Time	Source	Destination	Protocol	Length	Info
6	2021-07-23 12:50:40.896328907	10.9.0.5	192.168.60.5	UDP	48	55872 → 8080 Len=6
9	2021-07-23 12:50:50.861233367	10.9.0.5	192.168.60.6	UDP	48	55872 → 8080 Len=6
9	2021-07-23 12:50:55.327399046	10.9.0.5	192.168.60.7	UDP	48	55870 → 8080 Len=6
10	2021-07-23 12:50:57.621894978	10.9.0.5	192.168.60.5	UDP	48	52223 → 8080 Len=6

Using the random mode

在路由器上利用iptables命令，采用random模式创建负载均衡规则如下。

```
root@eaf163c47638:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic
--mode random --probability 0.33 -j DNAT --to-destination 192.168.60.5:8080
root@eaf163c47638:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic
--mode random --probability 0.5 -j DNAT --to-destination 192.168.60.6:8080
root@eaf163c47638:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -j DNAT --to-
destination 192.168.60.7:8080
```

在用户主机上向路由器的8080端口发送UDP数据包如下。

```
root@36cddb526b95:/# echo hello | nc -u 10.9.0.11 8080
^C
root@36cddb526b95:/# echo hello | nc -u 10.9.0.11 8080
^C
root@36cddb526b95:/# echo hello | nc -u 10.9.0.11 8080
^C
root@36cddb526b95:/# echo hello | nc -u 10.9.0.11 8080
^C
root@36cddb526b95:/# echo hello | nc -u 10.9.0.11 8080
^C
root@36cddb526b95:/# echo hello | nc -u 10.9.0.11 8080
^C
root@36cddb526b95:/# echo hello | nc -u 10.9.0.11 8080
^C
root@36cddb526b95:/# echo hello | nc -u 10.9.0.11 8080
^C
root@36cddb526b95:/# echo hello | nc -u 10.9.0.11 8080
^C
root@36cddb526b95:/# echo hello | nc -u 10.9.0.11 8080
^C
```

在服务器192.168.60.5上监听8080端口，得到结果如下。

```
root@44ca35555561:/# nc -luk 8080
hello
hello
hello
```

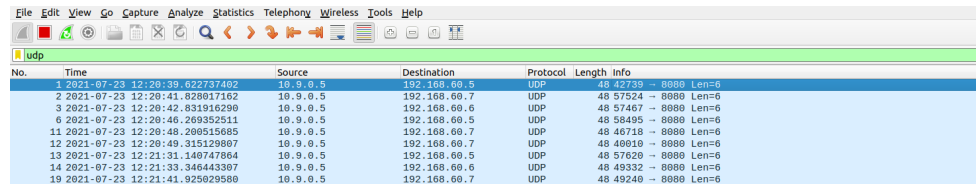
在服务器192.168.60.6上监听8080端口，得到结果如下。

```
root@5479b7208878:/# nc -luk 8080
hello
hello
```

在服务器192.168.60.7上监听8080端口，得到结果如下。

```
root@df14eb7bfad1:/# nc -luk 8080
hello
hello
hello
hello
```

利用wireshark抓包，得到结果如下，可知UDP数据包发送成功，且根据负载均衡规则随机发给三个服务器。



The image shows a Wireshark packet capture window with the 'udp' filter selected. The packet list pane displays 19 captured packets. The packet details pane shows the structure of a UDP packet, including Ethernet II, Internet Protocol Version 4, and User Datagram Protocol. The packet bytes pane shows the raw data of the selected packet.

No.	Time	Source	Destination	Protocol	Length	Info
1	2021-07-23 12:20:39.622737482	10.9.0.5	192.168.60.5	UDP	48	42735 → 8080 Len=6
2	2021-07-23 12:20:41.828017162	10.9.0.5	192.168.60.7	UDP	48	57524 → 8080 Len=6
3	2021-07-23 12:20:42.831916290	10.9.0.5	192.168.60.6	UDP	48	57467 → 8080 Len=6
6	2021-07-23 12:20:46.269352511	10.9.0.5	192.168.60.5	UDP	48	56495 → 8080 Len=6
11	2021-07-23 12:20:48.288515685	10.9.0.5	192.168.60.7	UDP	48	46718 → 8080 Len=6
12	2021-07-23 12:20:49.315129887	10.9.0.5	192.168.60.7	UDP	48	46819 → 8080 Len=6
13	2021-07-23 12:21:31.140747864	10.9.0.5	192.168.60.5	UDP	48	57620 → 8080 Len=6
14	2021-07-23 12:21:33.346443387	10.9.0.5	192.168.60.6	UDP	48	49332 → 8080 Len=6
19	2021-07-23 12:21:41.925029580	10.9.0.5	192.168.60.7	UDP	48	49240 → 8080 Len=6